



**UNIVERSITÀ
DEGLI STUDI
DI UDINE**
hic sunt futura

Usage of ClusterR Package

Alessio Corrado
Francesco Zuccato

2022/2023

Applied Statistics & Data Analysis

Department of Mathematics, Computer Science and Physics

University of Udine

Contents


- Introduction
 - Unsupervised learning
 - Cluster analysis
 - Data generation
- ClusterR Package
 - Provided models for clustering
 - Execution times
 - Centroid vs medoids
 - Provided datasets

Unsupervised Learning

- Cluster Analysis methods
 - Connectivity based
 - Agglomerative (Divisive) Hierarchical Clustering
 - Centroid based
 - K-Means, K-Medoids
 - Distribution based
 - Gaussian Mixture Models
 - Density based
 - DBSCAN


1. Connectivity based

- Create connections between clusters
- Agglomerative (Divisive) Hierarchical Clustering



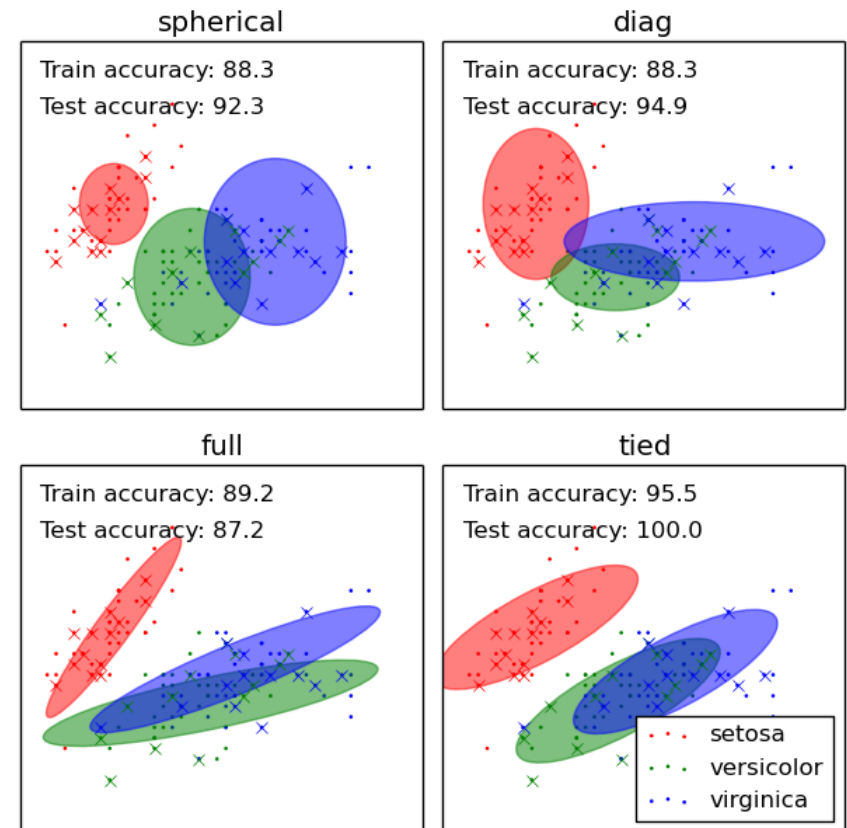
2. Centroid based

- Centroid (K-means)
- Medoid (K-medoid)




3. Distribution based

- Data explained by a probabilistic model
 - each observation belongs to the cluster that maximizes the generating probability
- Generalizes centroid-based
 - all centroids has the same gaussian distribution
- Gaussian Mixture Models



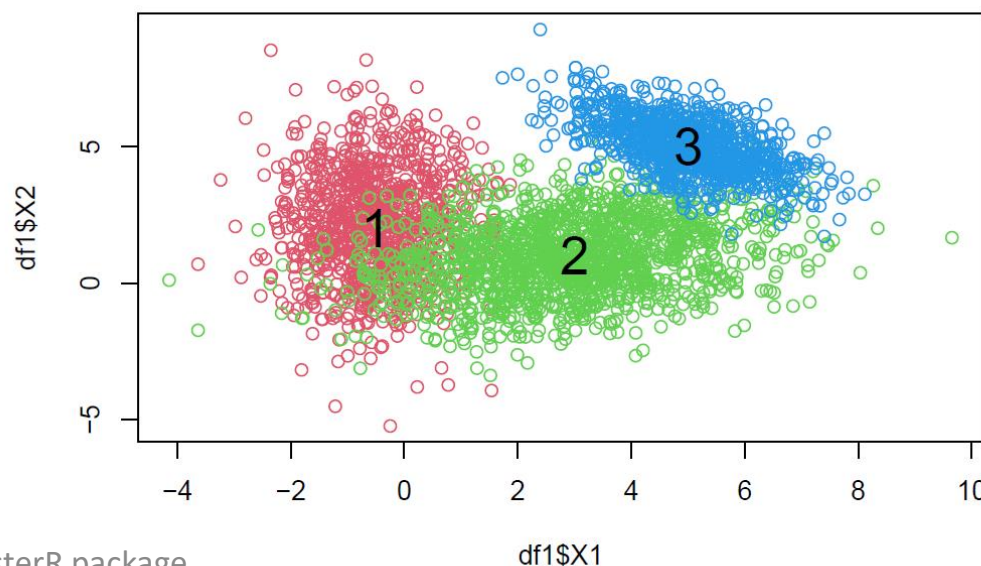

4. Density based

- Distance between points
- Dense connected regions
- No prior number of clusters
- DBSCAN, OPTICS



Data Generation

- Synthetic dataset
 - 3 clusters
 - 3 gaussian distributions
- *mvrnorm* R function
 - MASS package
 - n : number samples
 - μ : mean array
 - Σ : covariance matrix



ClusterR Package

- Clustering Methods
 - Gaussian Mixture Models: *GMM*
 - K-Means: *Kmeans_arma*, *Kmeans_rcpp*
 - K-Medoids: *Cluster_Medoids*, *Clara_Medoids*
- Automated choice of K
- Other Methods
 - Plot 2D
 - Plot Silhouette Dissimilarity
 - External Validation

1. Gaussian Mixture Models

Input


- (mandatory) **data**: the matrix with the observations (one row per item, one column per component)
- (mandatory) **n_gaus**: the number of gaussian processes
- **dist_mode**: specifies if the training algorithm should use euclidean or manhattan distance
- **seed_mode**: specifies the initial placement of the centroids for the iterative algorithm (static/random subset/spread)
- **km_iter**: number of iterations for the k-means algorithm
- **em_iter**: number of iterations for the expectation-maximization algorithm
- **var_floor**: smallest possible values for diagonal covariances
- **seed**: integer for the random number generator (to allow replicability)

```
gmm_3 = GMM(df1_no_label, 3, dist_mode = "maha_dist", seed_mode = "random_subset")
gmm_3 = reorder_data(gmm_3, "centroids", list("weights", "covariance_matrices"))
df1["y_3gmm"] = predict(gmm_3, newdata = df1_no_label)
```

1. Gaussian Mixture Models

Output

- **centroids:** a matrix $\in \mathbb{R}^{k \times d}$ that specifies the position of each centroid
- **covariance_matrices:** a matrix $\in \mathbb{R}^{k \times d}$ that specifies the diagonal values of each covariance
- **weights:** a vector $\in \mathbb{R}^k$, with the percentage of weight of each gaussian component
- **Log_likelihood:** a matrix $\in \mathbb{R}^{n \times k}$ foreach training item
- **call:** a list containing the values of the parameters used in the invocation



```
gmm_3$centroids
```

```
##           [,1]      [,2]
## [1,] -0.4047667  1.9044835
## [2,]  3.0055130  0.8478915
## [3,]  4.9509768  4.7739142
```

```
gmm_3$covariance_matrices
```

```
##           [,1]      [,2]
## [1,]  0.8987419  4.099430
## [2,]  2.4693921  1.496995
## [3,]  1.0732790  1.408025
```

```
gmm_3$weights
```

```
## [1]  0.3183000  0.3538941  0.3278059
```

```
head(gmm_3$Log_likelihood)
```

```
##           [,1]      [,2]      [,3]
## [1,] -21.07133 -8.399115 -3.217083
## [2,] -31.93476 -28.100360 -9.977866
## [3,] -14.78147 -4.384024 -2.784711
## [4,] -19.08291 -4.379669 -2.689926
## [5,] -21.52626 -4.488746 -4.122479
## [6,] -10.85150 -3.306210 -4.393683
```

2. K-means

Arma vs Rcpp

Shared parameters:

- (mandatory) **data**: matrix with the observations
- (mandatory) **clusters**: number of clusters
- **CENTROIDS**: matrix with the initial cluster centroids
- **verbose**: whether or not to print some logs during the process

KMeans_arma pros:

- *ready-to-go approach*
- faster than *KMeans_rcpp*:
 - returns only the matrix of the centroids
 - can work in parallel with OpenMP


KMeans_rcpp pros:

- Many more initialization and running settings

2. K-means

Arma: *Armadillo* Implementation

```
kmeans_arma = list()
kmeans_arma$centroids = matrix(KMeans_arma(df1_no_label, 3), nrow=3)
kmeans_arma = reorder_data(kmeans_arma, "centroids", list())
df1["y_3Mn_arma"] = predict_KMeans(df1_no_label, kmeans_arma$centroids)
```



2. K-means

Rcpp: *RcppArmadillo* Implementation

Additional input parameters:

- initialize various parameters such as **initializer**, **fuzzy**, **tol_optimal_init**, **seed**
- set the running time and convergence:
 - **num_init**: number of different initializations, if > 1 then is returned the best fit, according to within-cluster-sum-of-squared-error (WCSS)
 - **max_iters**, **tol**: stopping criterias based on the number of iterations or the accuracy


```
kmeans_rcpp = KMeans_rcpp(df1_no_label, 3)
kmeans_rcpp = reorder_data(kmeans_rcpp, "centroids", list("WCSS_per_cluster", "obs_per_cluster"))
kmeans_rcpp$clusters <- NULL #drop it now, it's wrong, use the predict below
df1["y_3Mn_rcpp"] = predict(kmeans_rcpp, newdata = df1_no_label)
```


2. K-means

Rcpp: *RcppArmadillo* Implementation

- Output:

- **call**, **centroids**: as for the GMM
- **clusters**: a vector of n elements with the predicted cluster for each item
- **best_initialization**: an integer indicating which was the best initialization (useful when $num_init > 1$)
- some metrics: **total_SSE** (squared distance of each point to its centroid), **WCSS_per_cluster**, **obs_per_cluster** (counts items predicted per cluster)



```
> kmeans_rcpp$total_SSE
[1] 39846.24
> kmeans_rcpp$best_initialization
[1] 1
> kmeans_rcpp$WCSS_per_cluster
[1] 5627.612 3098.745 3445.179
> kmeans_rcpp$obs_per_cluster
[1] 1229 1161 1110
> kmeans_rcpp$between_SS_DIV_total_SS
[1] 0.6945374
```

3. K-medoids

Partitioning Around Medoids (PAM)

1. BUILD

- I. Choose first medoid: the most central object
- II. Choose (k-1) medoids: objects that further minimize the overall dissimilarity

2. SWAP

- I. Foreach pair of (x, y) where x is a medoid and y isn't:
 - I. Calculate objective cost (distance between any pair of points)
 - II. If it is decreased $SWAP(x, y)$

Note: since it computes the dissimilarity matrix M for any pair of elements it's trivially a **quadratic** algorithm.

3. K-medoids

Cluster vs Clara

- *Cluster_Medoids*
 - Computes M for any couple (x, y)
 - Finds an exact solution
 - Ideal for small datasets
- *Clara_Medoids*
 - **Clustering *LAR*ge *A*pplications**
 - Uses only a sampled subset of the whole dataset
 - Computes M only for the sampled elements
 - Finds an approximated solution
 - Ideal for big datasets

3. K-medoids

Cluster vs Clara

Shared parameters:

- (mandatory) **data**, **clusters**: as for the k-means
- **distance_metric**: the distance method to be used: euclidean, manhattan, chebyshev, hamming, etc.
- **threads**: number of cores (for parallelism)
- **swap_phase**: whether or not to apply also the SWAP phase

Additional Clara parameters:

- (mandatory) **samples**: number of samples to draw from the data set
- (mandatory) **sample_size** $\in (0, 1]$: percentage of the data to draw in each sample iteration

3. K-medoids

Cluster vs Clara

Output (of Cluster):

- `call`, `medoids`, `clusters`: as for the `Kmeans` (with medoids instead of centroids)
- `silhouette_matrix`: dataframe $\in \mathbb{R}^{n \times 7}$
- `dissimilarity_matrix`: matrix $\in \mathbb{R}^{n \times n}$ with the distance foreach couple of items
- `best_dissimilarity`: an overall statistics
- `clustering_stats` dataframe $\in \mathbb{R}^{k \times 6}$ with per-cluster statistics: `clusters`, `average_dissimilarity`, `max_dissimilarity`, `diameter`, `separation`, `number_obs`

For Clara few differences:


- Smaller M (only for the sampled)
- Removed: `diameter` and `separation`
- Added: `isolation`

3. K-medoids

Cluster vs Clara


Cluster Medoids

```
kmedoids_pam = Cluster_Medoids(df1_no_label, 3, distance_metric="euclidean")
```



Clara Medoids

```
kmedoids_cla = Clara_Medoids(df1_no_label, 3, samples=4, sample_size=0.1, "euclidean")
```



Finding optimal number of clusters

- *Optimal_Clusters_[method]*
- [Method]: GMM, KMeans, KMedoids

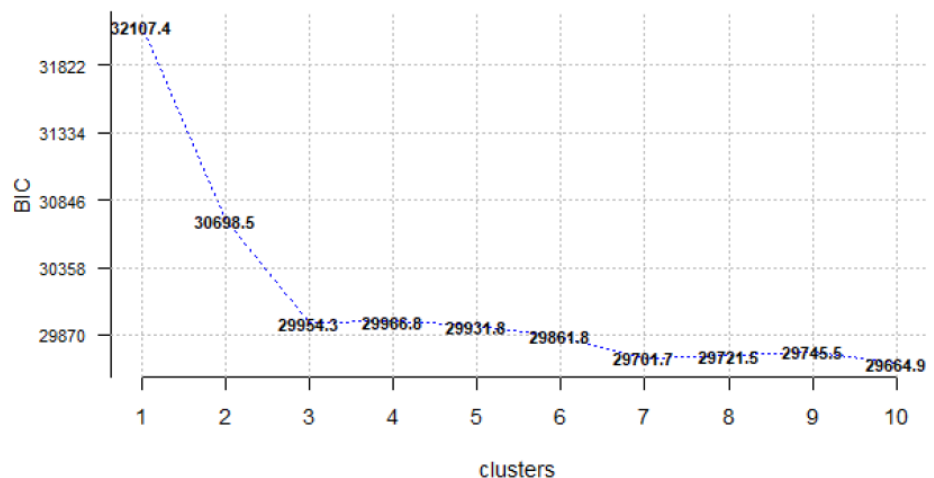
```
opt_gmm = Optimal_Clusters_GMM(df1_no_label, max_clusters = 10, criterion = "BIC",  
                                dist_mode = "eucl_dist", seed_mode = "random_subset", plot_data = T)
```

Additional parameters:

- max_clusters, criterion

For the choice (simplest way):


1. Choose a criterion
2. Plot: k vs criterion
3. Choose k: Elbow method



Plot_2d, External_validation

- *Plot_2d*: generates a scatter plot which colors the clusters, adds the centroids and the legend. Params:
 - Clusters: predictions
 - Centroids_medoids
- *External_validation*: given the real clusters and the predicted, prints some metrics. Params:
 - real, predicted clusters
 - method: to return, default adj R²

```
plot_2d(data = df1_no_label, clusters = df1$y_3gmm, centroids_medoids = gmm_3$centroids)
```



```
external_validation(df1$color, df1$y_3gmm, summary_stats = T)
```

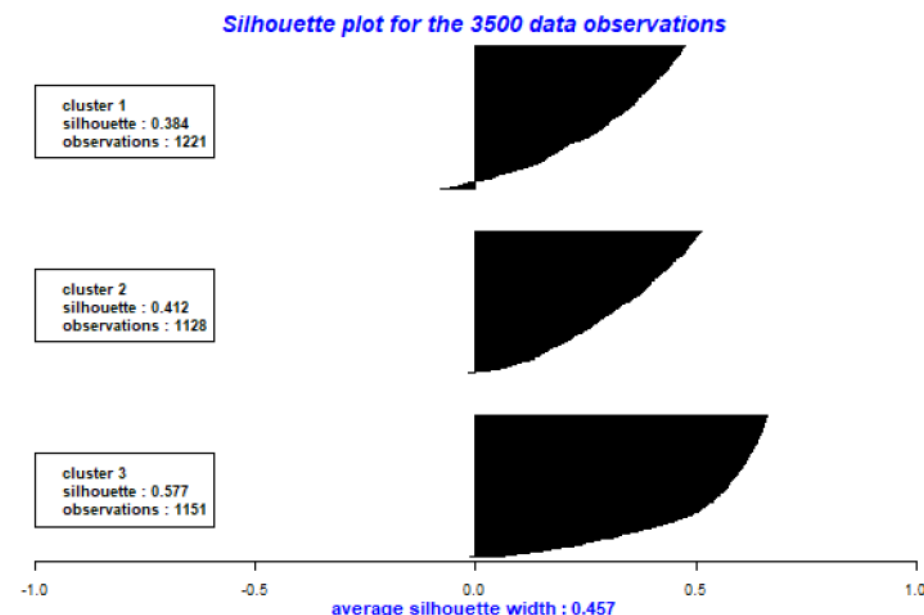
```
##
## -----
## purity                : 0.898
## entropy                : 0.3053
## normalized mutual information : 0.7007
## variation of information : 0.9402
## normalized var. of information : 0.4608
## -----
## specificity            : 0.9107
## sensitivity            : 0.7934
## precision              : 0.825
## recall                 : 0.7934
## F-measure              : 0.8089
## -----
## accuracy OR rand-index : 0.87
## adjusted-rand-index     : 0.7104
## jaccard-index           : 0.6791
## fowlkes-mallows-index   : 0.809
```

Silhouette Dissimilarity Plot

Average Silhouette Width

- Visual way to evaluate the clusters built with the medoids algorithm
- **Cohesion:** object should be similar to others in its own cluster
- **Separation:** object should be dissimilar to others in other clusters
- Range: $[-1, +1]$
- Goal: to maximize it

```
invisible(Silhouette_Dissimilarity_Plot(kmedoids_pam, silhouette = TRUE))
```

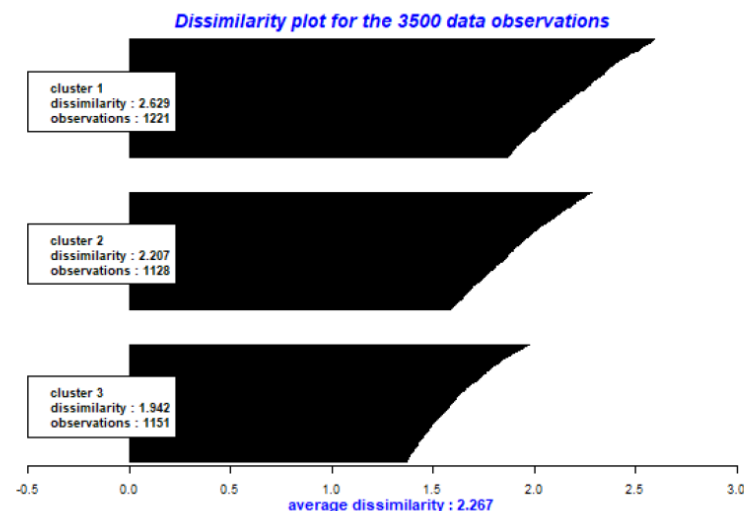


Silhouette Dissimilarity Plot

Average Dissimilarity

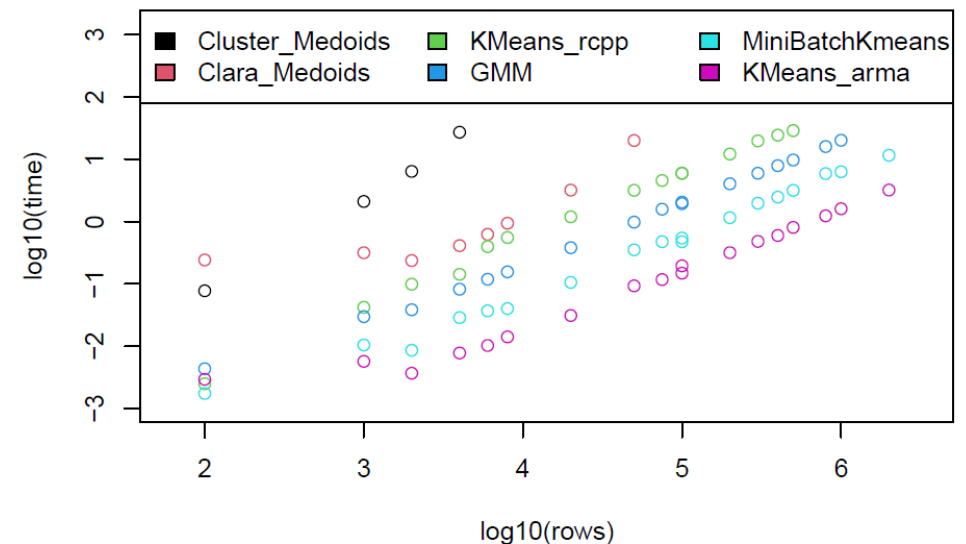
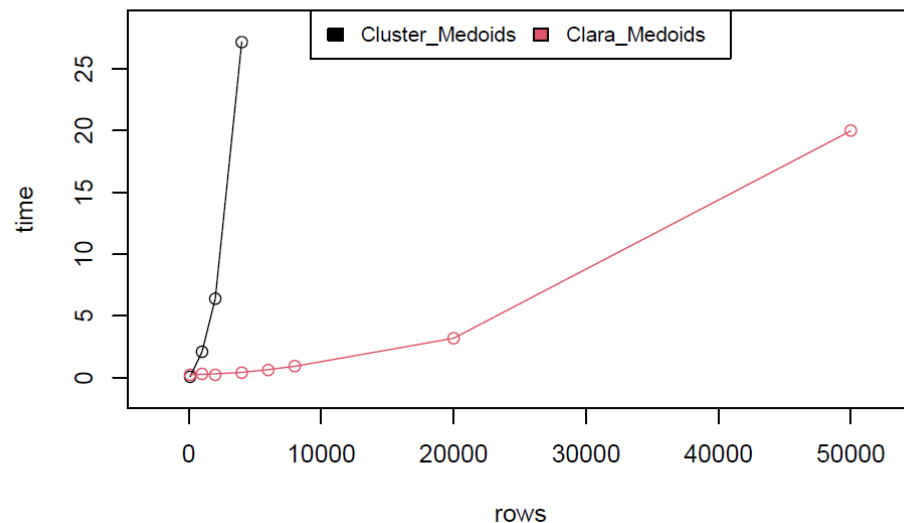

- Visual way to evaluate the clusters built with the medoids algorithm
- **Dissimilarity**: calculated per cluster

```
invisible(Silhouette_Dissimilarity_Plot(kmedoids_pam, silhouette = FALSE))
```



Execution times

- We compared the execution time of the various algorithms
 - n = #observations: from 100 to 2 mln
 - p = #dimensions: 10
 - Data generated with a uniform distribution on $[0,1]$ with $n * p$ samples



Execution times

Considerations

- For each algorithm with *drop1* we compared three different linear models:
 - **L**: Linear
 - **Q**: Quadratic
 - **LQ**: Linear + Quadratic
- Results
 - K-Means, GMM fit the **L** model: are linear
 - K-Medoids fit the **Q** model: are quadratic

Cluster Medoids

Execution Times

- The AIC slightly suggests Q vs LQ model
- The Q model seems fair
 - The quadratic effect is confirmed by p-value

```
t = df_times %>% filter(method == 'Cluster_Medoids')
model_l = lm(time ~ rows, data = t, weights=1/rows)
model_q = lm(time ~ rows_squared, data = t, weights=1/rows)
model_lq = lm(time ~ rows + I(rows_squared), data = t, weights=1/rows)
drop1(model_lq) # drop1: linear vs quadratic vs both
```

```
## Single term deletions
##
## Model:
## time ~ rows + I(rows_squared)
##
```

	Df	Sum of Sq	RSS	AIC
<none>			0.0002106	-33.407
rows	1	0.0000012	0.0002118	-35.384
I(rows_squared)	1	0.0237289	0.0239396	-16.474

```
summary(model_q) # summary of the model with lowest BIC
##
## Call:
## lm(formula = time ~ rows_squared, data = t, weights = 1/rows)
##
## Weighted Residuals:
##      1      2      3      4
## -0.001616  0.010783 -0.009433  0.001993
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.686e-02  9.786e-02   0.785 0.514510
## rows_squared 1.685e-06  3.926e-08  42.922 0.000542 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01029 on 2 degrees of freedom
## Multiple R-squared:  0.9989, Adjusted R-squared:  0.9984
## F-statistic: 1842 on 1 and 2 DF, p-value: 0.0005424
```

Clara Medoids

Execution Times

- Same considerations apply for Clara, of course the quadratic effect depends only the number of sampled elements
- The BIC slightly suggests Q vs LQ model
- The Q model seems fair
 - The quadratic effect is confirmed by p-value

```
t = df_times %>% filter(method == 'Clara_Medoids')
model_l = lm(time ~ rows, data = t, weights=1/rows)
model_q = lm(time ~ rows_squared, data = t, weights=1/rows)
model_lq = lm(time ~ rows + I(rows_squared), data = t, weights=1/rows)
drop1(model_lq, k=log(length(df_times$rows))) #drop1: linear vs quadratic
```

```
## Single term deletions
##
## Model:
## time ~ rows + I(rows_squared)
##              Df Sum of Sq      RSS      AIC
## <none>                  0.00001198 -94.149
## rows                1 0.00000118 0.00001316 -97.777
## I(rows_squared)    1 0.00184281 0.00185479 -58.191
```

```
summary(model_q) # summary of the model with lowest BIC
```

```
##
## Call:
## lm(formula = time ~ rows_squared, data = t, weights = 1/rows)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0014571 -0.0008033  0.0003704  0.0013373  0.0020965
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.499e-01  1.349e-02  18.52 1.60e-06 ***
## rows_squared  7.886e-09  1.284e-10  61.43 1.25e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.001481 on 6 degrees of freedom
## Multiple R-squared:  0.9984, Adjusted R-squared:  0.9981
## F-statistic: 3774 on 1 and 6 DF, p-value: 1.25e-09
```

K-Means Rcpp

Execution Times

- The BIC slightly suggests L vs LQ model
- The L model seems fair
 - The linear effect is confirmed by p-value
 - In any case, the quadratic coefficient in the LQ model is insignificant (and negative!)
- For Arma and MiniBatch hold same considerations

```
model_lq$coefficients # check coefficient values of LQ
```

```
##      (Intercept)          rows I(rows_squared)
## -7.858108e-03    6.291618e-05   -6.660551e-12
```

```
t = df_times %>% filter(method == 'KMeans_rcpp')
model_l = lm(time ~ rows, data = t, weights=1/rows)
model_q = lm(time ~ rows_squared, data = t, weights=1/rows)
model_lq = lm(time ~ rows + I(rows_squared), data = t, weights=1/rows)
drop1(model_lq, k=log(length(df_times$rows))) #drop1: linear vs quadratic
```

```
## Single term deletions
##
## Model:
## time ~ rows + I(rows_squared)
##              Df Sum of Sq      RSS      AIC
## <none>                  0.00001439 -194.71
## rows                1 0.00130480 0.00131919 -131.32
## I(rows_squared)     1 0.00000189 0.00001628 -197.24
```

```
summary(model_l) # summary of the model with lowest BIC
```

```
##
## Call:
## lm(formula = time ~ rows, data = t, weights = 1/rows)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -2.165e-03 -3.750e-04 -1.949e-05  3.800e-04  2.714e-03
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.162e-03  1.020e-02  -0.604    0.556
## rows        6.076e-05  8.465e-07  71.772 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.001119 on 13 degrees of freedom
## Multiple R-squared:  0.9975, Adjusted R-squared:  0.9973
## F-statistic: 5151 on 1 and 13 DF, p-value: < 2.2e-16
```

GMM

Execution Times

- The BIC slightly suggests L vs LQ model
- The L model seems fair
 - The linear effect is confirmed by p-value
 - In any case, the quadratic coefficient in the LQ model is insignificant and has a very high p-value

```
t = df_times %>% filter(method == 'GMM')
model_l = lm(time ~ rows, data = t, weights=1/rows)
model_q = lm(time ~ rows_squared, data = t, weights=1/rows)
model_lq = lm(time ~ rows + I(rows_squared), data = t, weights=1/rows)
drop1(model_lq, k=log(length(df_times$rows))) #drop1: linear vs quadr

## Single term deletions
##
## Model:
## time ~ rows + I(rows_squared)
##
```

	Df	Sum of Sq	RSS	AIC
<none>			4.1500e-07	-284.85
rows	1	2.92e-04	2.9242e-04	-177.74
I(rows_squared)	1	3.30e-08	4.4700e-07	-287.94

```
summary(model_l) # summary of the model with lowest BIC
##
## Call:
## lm(formula = time ~ rows, data = t, weights = 1/rows)
##
## Weighted Residuals:
```

	Min	1Q	Median	3Q	Max
	-3.463e-04	-9.100e-05	-3.702e-05	1.135e-04	2.973e-04

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.673e-03	1.571e-03	1.701	0.11
rows	1.995e-05	9.176e-08	217.436	<2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0001727 on 15 degrees of freedom
## Multiple R-squared:  0.9997, Adjusted R-squared:  0.9997
## F-statistic: 4.728e+04 on 1 and 15 DF, p-value: < 2.2e-16
```


Centroids vs Medoids

- K-Medoids pros: robustness to noise and outliers
- K-Means pros: fast (linear)
- We gave an example of the two algorithms with 10 representations


The Outlier Effect



K-Means



K-Medoids




Dataset 1: Dietary Survey IBS

- Synthetic data
- 400 rows, 42 explanatory variables
- 1 binary response variable *class*:
 - healthy or IBS-positive (50% each)
- Performed PCA and run KMeans_rcpp


```
colnames(dietary_survey_IBS)
```

```
## [1] "bread"           "wheat"
## [3] "pasta"           "breakfast_cereal"
## [5] "yeast"           "spicy_food"
## [7] "curry"           "chinese_takeaway"
## [9] "chilli"          "cabbage"
## [11] "onion"           "garlic"
## [13] "potatoes"        "pepper"
## [15] "vegetables_unspecified" "tomato"
## [17] "beans_and_pulses" "mushroom"
## [19] "fatty_foods_unspecified" "sauces"
## [21] "chocolate"       "fries"
## [23] "crisps"           "desserts"
## [25] "eggs"            "red_meat"
## [27] "processed_meat"   "pork"
## [29] "chicken"          "fish_shellfish"
## [31] "dairy_products_unspecified" "cheese"
## [33] "cream"           "milk"
## [35] "fruit_unspecified" "nuts_and_seeds"
## [37] "orange"          "apple"
## [39] "banana"          "grapes"
## [41] "alcohol"         "caffeine"
## [43] "class"
```

First two PCA scores, real classes



First two PCA scores, predicted classes



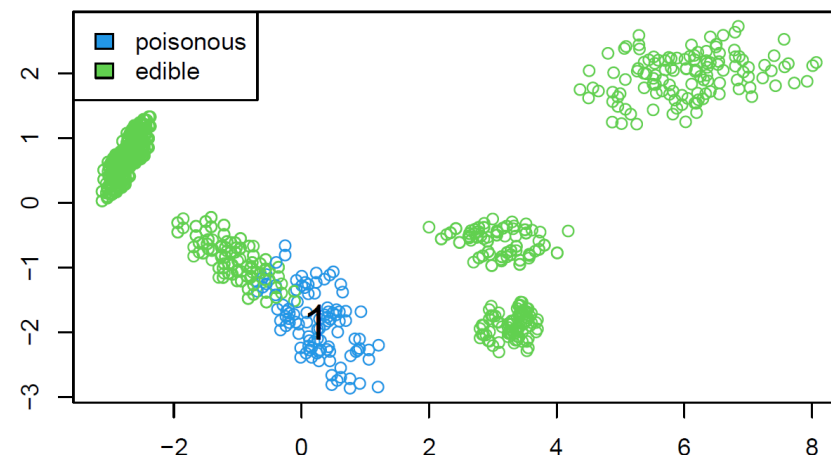
Dataset 2: Mushrooms

- Hypothetical samples of 23 species of gilled mushrooms
- 8124 rows, 22 explanatory variables
- 1 binary response variable *class*:
 - edible (48%) or poisonous (52%)
- Performed PCA and run KMeans_rcpp


```
colnames(mushroom)
```

```
## [1] "class"           "cap_shape"
## [3] "cap_surface"     "cap_color"
## [5] "bruises"         "odor"
## [7] "gill_attachment" "gill_spacing"
## [9] "gill_size"       "gill_color"
## [11] "stalk_shape"     "stalk_root"
## [13] "stalk_surface_above_ring" "stalk_surface_below_ring"
## [15] "stalk_color_above_ring" "stalk_color_below_ring"
## [17] "veil_type"       "veil_color"
## [19] "ring_number"     "ring_type"
## [21] "spore_print_color" "population"
## [23] "habitat"
```

First two PCA scores, real classes




First two PCA scores, predicted classes



Dataset 3: Soybean

- Soybeans characteristics from the UCI machine learning repository
- 307 rows, 35 explanatory variables
- 1 categorical response variable *class*:
 - 19 different classes
- Performed PCA and ...
 - Classes not separable in 2D

First two PCA scores, real classes



```
colnames(soybean)
```

```
## [1] "date"           "plant_stand"    "precip"         "temp"
## [5] "hail"           "crop_hist"      "area_damaged"   "severity"
## [9] "seed_tmt"       "germination"    "plant_growth"   "leaves"
## [13] "leafspots_halo" "leafspots_marg" "leafspot_size"  "leaf_shread"
## [17] "leaf_malf"      "leaf_mild"      "stem"           "lodging"
## [21] "stem_cankers"   "canker_lesion"  "fruiting_bodies" "external_decay"
## [25] "mycelium"       "int_discolor"   "sclerotia"      "fruit_pods"
## [29] "fruit_spots"    "seed"           "mold_growth"    "seed_discolor"
## [33] "seed_size"      "shriveling"     "roots"          "class"
```